

CN-Basic

L23


Reliable Data Transfer

Higher Version

Dr. Ram P Rustagi
rprustagi@ksit.edu.in
<http://www.rprustagi.com>
<https://www.youtube.com/rprustagi>

Chapter 3

Transport Layer



A note on the use of these ppt slides:

We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- ❖ If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- ❖ If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

Thanks and enjoy! JFK/KWR

All material copyright 1996-2012
J.F Kurose and K.W. Ross, All Rights Reserved



*Computer
Networking: A Top
Down Approach*
6th edition
Jim Kurose, Keith Ross
Addison-Wesley
March 2012


rdt so far

- **rdt 1 . 0**
 - Fully reliable underlying channel
- **rdt 2 . 0**
 - No packet loss, but packet corruption can occur
 - Receiver needs to send ACK/NAK
- **rdt 2 . 1**
 - Introduced seq number to deal with duplicate pkts
- **rdt 2 . 2**
 - Used ACK instead of NAK
- **rdt 3 . 0**
 - Underlying channel can also lose packets
 - Timer required for retransmit
 - Efficiency poor for stop-and-wait protocols
- **Next ?**


Pipelined protocols

Pipelining: sender sends multiple, “in-flight”, to-be-acked

- Range of sequence numbers must be increased
- Buffering at sender and/or receiver
- Increases efficiency




(a) a stop-and-wait protocol in operation



(b) a pipelined protocol in operation

- Two generic forms of pipelined protocols:
 - *Go-Back-N (GBN), Selective Repeat(SR)*

Pipelining: Increased Utilization



3-packet pipelining increases utilization by a factor of 3!

$$U_{\text{sender}} = \frac{3L / R}{RTT + L / R} = \frac{.0024}{30.008} = 0.00081$$


Pipelined protocols: overview

Key point: sender can have up to N unacked pkts in pipeline

- Go-back-N:
- Receiver only sends *cumulative ack*
 - Doesn't ack packet if there's a gap
 - Acks last in-seq pkt received
- Sender has timer for oldest unacked packet
 - When timer expires, retransmit *all* unacked packets
- Selective Repeat:
- Rcvr sends *individual ack* for each packet
 - Buffer requirement?
- Sender maintains timer for each unacked packet
- When timer expires, retransmit only that unacked packet

Go-Back-N: sender

- k-bit seq # in pkt header
- “window” of up to N, consecutive unack’ed pkts allowed



- ACK(n): ACKs all pkts up to, including seq # n - *“cumulative ACK”*
- may receive duplicate ACKs (see receiver)
- timer for oldest in-flight pkt
- *timeout(n)*: retransmit packet n and all higher seq # pkts in window
- it is a **Sliding Window Protocol**.

Go-Back-N: sender

- Why to limit the window size to N
 - What happens if it is made larger (unlimited)?
 - Issues w.r.t flow control (study later)
 - Issues w.r.t. congestion control (study later)
- Given k bits for sequence number,
 - Sequence numbers are
 - $0, 1, 2, \dots, 2^k - 1, 0, \dots$
 - Need to use modulo 2^k arithmetic
 - Roll over of seq number has its own set of problems
- GBN Simulation Applet
 - http://www.ccs-labs.org/teaching/rn/animations/gbn_sr/
 - <http://computerscience.unicam.it/marcantoni/reti/applet/GoBackProtocol/goback.html>

GBN protocol: example

- Consider the following case
 - Window size = 4
 - Sequence number 0, 1, ..., 6 (no roll over)
 - Pkt with seq=2 (i.e. 3rd) is lost and hence times out
 - A total of 6 packets are to be transmitted
- Exercise:
 - Draw the timeline of packets and ack communication

GBN in action

sender window (N=4)

0 1 2 3 4 5 6 7 8
 0 1 2 3 4 5 6 7 8
 0 1 2 3 4 5 6 7 8
 0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8
 0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8
 0 1 2 3 4 5 6 7 8
 0 1 2 3 4 5 6 7 8
 0 1 2 3 4 5 6 7 8

sender

send pkt0
 send pkt1
 send pkt2
 send pkt3
 (wait)

rcv ack0, send pkt4
 rcv ack1, send pkt5
 ignore duplicate ACK



pkt 2 timeout

send pkt2
 send pkt3
 send pkt4
 send pkt5

receiver

receive pkt0, send ack0
 receive pkt1, send ack1

receive pkt3, discard,
 (re)send ack1

receive pkt4, discard,
 (re)send ack1

receive pkt5, discard,
 (re)send ack1

rcv pkt2, deliver, send ack2
 rcv pkt3, deliver, send ack3
 rcv pkt4, deliver, send ack4
 rcv pkt5, deliver, send ack5

GBN: Sender Extended FSM

rdt_send(data)

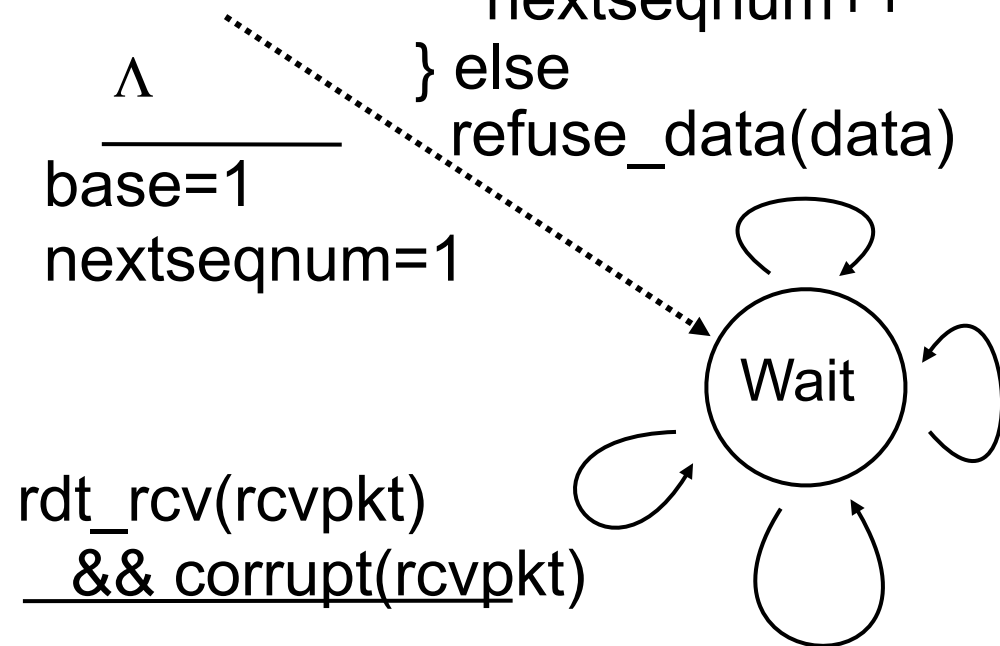
```

if (nextseqnum < base+N) { //ensure window is not full
    sndpkt[nextseqnum] = make_pkt(nextseq,data,chksum)
    udt_send(sndpkt[nextseqnum])
    if (base == nextseqnum)
        start_timer
    nextseqnum++
} else
    refuse_data(data)

```

Three event types

- invocation from above
- timeout
- receipt of an ack
 - corrupt or uncorrupt



timeout // gives the protocol name GBN

```

start_timer
udt_send(sndpkt[base])
udt_send(sndpkt[base+1])
...
udt_send(sndpkt[nextseqnum-1])

```

rdt_rcv(rcvpkt) && notcorrupt(rcvpkt)

base = getacknum(rcvpkt)+1 //cumulative ack mechanism

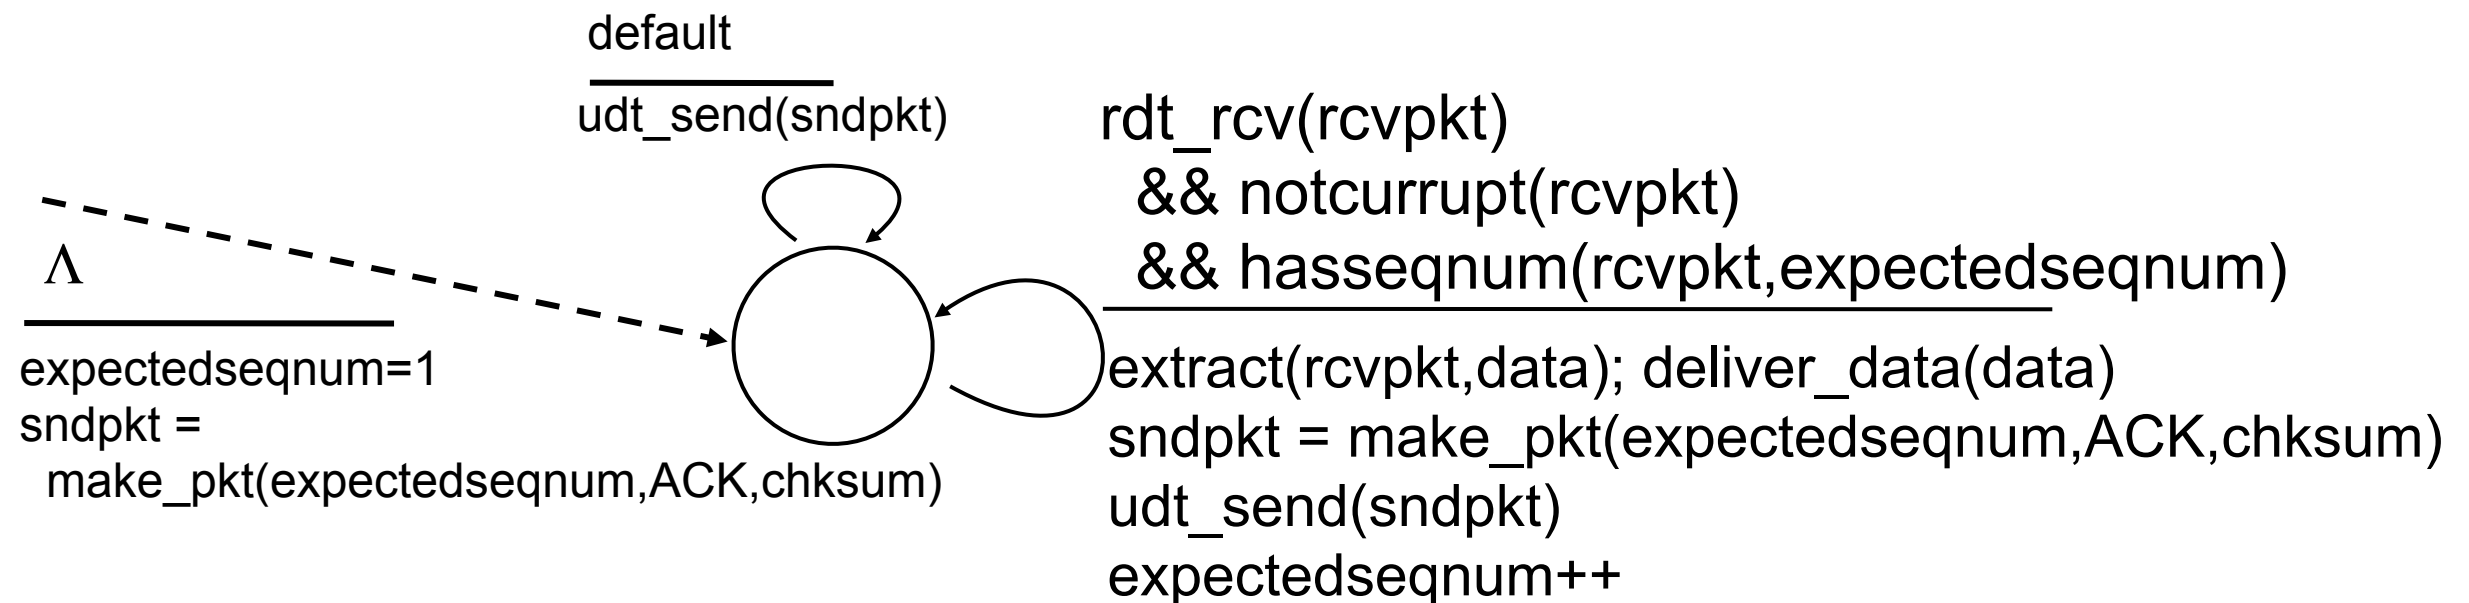
If (base == nextseqnum)

stop_timer

else

start_timer

GBN: receiver extended FSM



ACK-only: always send ACK for correctly-received pkt with highest *in-order* seq #

- May generate duplicate ACKs
- Need only remember **expectedseqnum**
- out-of-order pkt:
 - Discard (don't buffer): *no receiver buffering!*
 - *Does buffering help? Will sender anyway resend it?*
 - Re-ACK pkt with highest in-order seq #


GBN in action

- Benefits over stop-n-wait
 - Fills the pipeline, efficient channel utilization
- Issues
 - Performance problems if BW-delay product is large
 - A single packet error causes large num of retransmissions
 - Most of them unnecessary
 - With increasing error rates
 - Channel is filled with retransmissions
- Solution: **Selective Repeat**



Selective Repeat

- Receiver *individually* acknowledges all correctly received pkts
 - Buffers pkts, as needed, for eventual in-order delivery to upper layer
 - Acks need not be in order
- Sender only resends pkts for which ACK not received
 - Sender timer for each unACKed pkt
 - More complexity at sender
- Sender window
 - N consecutive seq #'s
 - Limits seq #'s of sent, unACKed pkts

Selective repeat: sender, receiver windows



Sender view of sequence number



Receiver view of sequence number

Selective repeat in action


- SR Simulation Applet
 - <http://computerscience.unicam.it/marcantoni/reti/applet/SelectiveRepeatProtocol/selRepProt.html>
- Observation
 - Receiver re-acknowledges packets even below its current window base
 - Is it really needed?
 - Sender and receiver windows typically do not coincide
 - Causes issues w.r.t. finite range of seq numbers

Selective repeat in action

- Consider the following scenario
 - Window size of 4
 - Seq number starting from 0, 1, ...
 - pkt0, pkt1, pkt2, pkt3, pkt4, pkt5, ...
 - pkt2 is lost
 - Draw the Timeline Sequence diagram for SR protocol

Selective repeat in action

sender window (N=4)



sender

send pkt0
 send pkt1
 send pkt2
 send pkt3
 (wait)

rcv ack0, send pkt4
 rcv ack1, send pkt5

record ack3 arrived



pkt 2 timeout

send pkt2

record ack4 arrived

record ack5 arrived

receiver

receive pkt0, send ack0
 receive pkt1, send ack1

receive pkt3, buffer,
 send ack3

receive pkt4, buffer,
 send ack4

receive pkt5, buffer,
 send ack5

rcv pkt2; deliver pkt2,
 pkt3, pkt4, pkt5; send ack2

Q: what happens when ack2 arrives?

Selective repeat(actions by sender/receiver)

sender

- Data from above:
 - If next available seq # in window, send pkt, start timer
- Timeout(n):
 - Resend pkt n, restart timer
- ACK(n) in [sendbase,sendbase+N]:
 - Mark pkt n as received
- If n smallest unACKed pkt, advance window base to next unACKed seq #

receiver


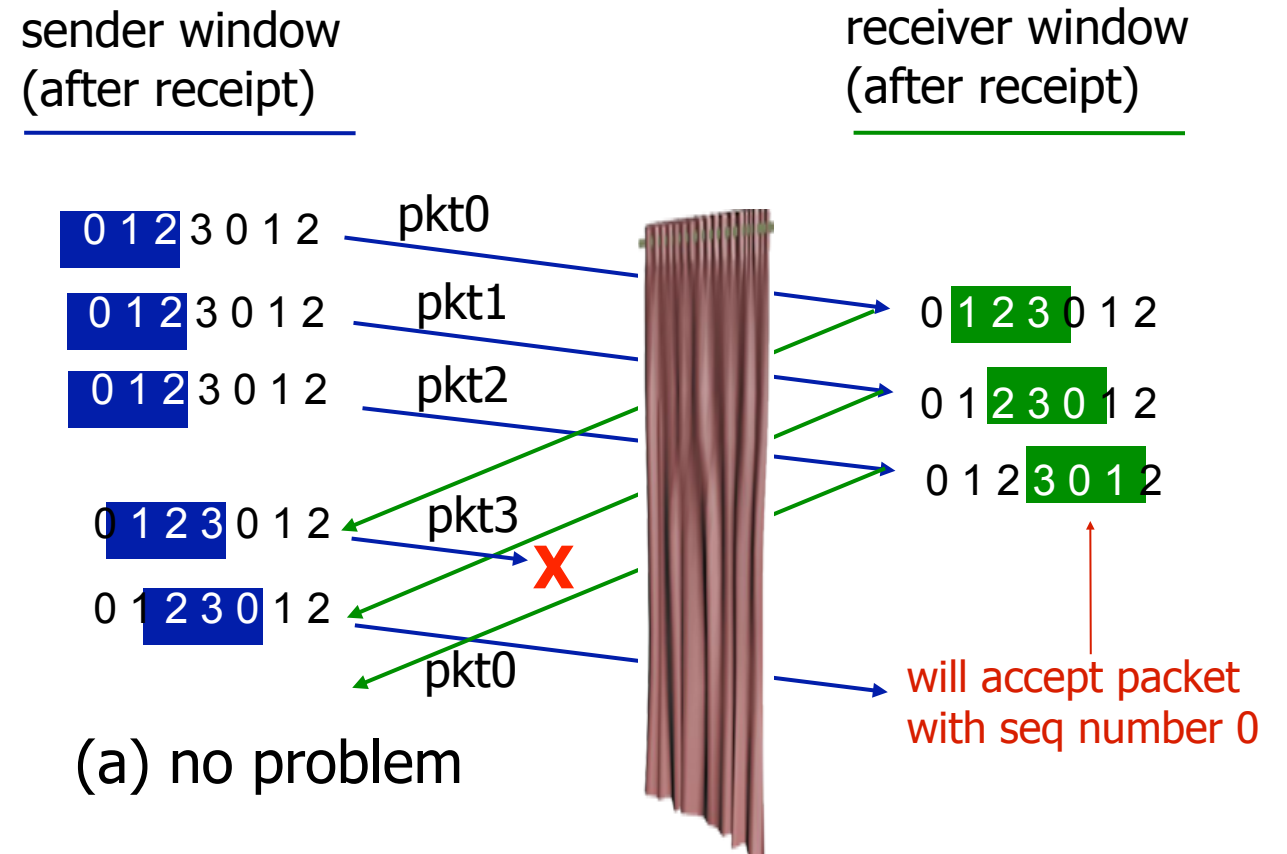
- Pkt n in [rcvbase, rcvbase+N-1]
 - Send ACK(n)
 - Out-of-order: buffer
 - In-order: deliver (also deliver buffered, in-order pkts), advance window to next not-yet-received pkt
- pkt n in [rcvbase-N,rcvbase-1]
 - ACK(n)
- Otherwise:
 - Ignore

Selective repeat in action

- Home Exercise
 - Draw the FSM for Sender and receiver
- Issues with SR
 - When seq numbers range is close to window size?
 - Window size 3,
 - Sequence number (only 4 values) 0, 1, 2, 3.

Selective Repeat: dilemma

- Example:
- Seq #'s: 0, 1, 2, 3
- Window size=3
- Receiver sees no difference in two scenarios!
- Duplicate data accepted as new in (b)
- Q: What relationship between seq # size and window size to avoid problem in (b)?



Selective repeat dilemma

- Figurative curtain between sender and receiver
 - Receiver can't see action taken by sender
 - From receiver perspective
 - Two scenario (prev slide) are identical
 - No way to distinguish between
 - Retransmission of pkt #0 with original pkt #5
- Issue:
 - Window size of 1 less than size of seq number
 - How small the window size should be?
 - Less than $N/2$

Summary of **rdt**

- Numerous mechanism for **rdt**(reliable data transfer)
 - Checksum, timer, seq number, ACK, NAK, window size, pipelining
- Mechanisms were incrementally added
 - To address increasingly complex (realistic) problems
- Underlying assumption
 - Channel sends packets in order, does not reorder them
 - Not true in real life
 - Multiple paths between sender and receiver
 - Impact of packet re-order
 - Pkt with old seq num x (or ack) can re-appear
 - Though neither side contains x
 - How to address such duplicate packets
 - Sender should use a seq #, only when sure no such pkts exists in the network
 - Assumes pkt has a lifetime (TCP assumes it to be 3 minutes)