

# CN-Advanced L35

## Distributed Hash Table

Dr. Ram P Rustagi  
rprustagi@ksit.edu.in  
<http://www.rprustagi.com>  
<https://www.youtube.com/rprustagi>

# Acknowledgements

## Chapter

## Multimedia

## Networking

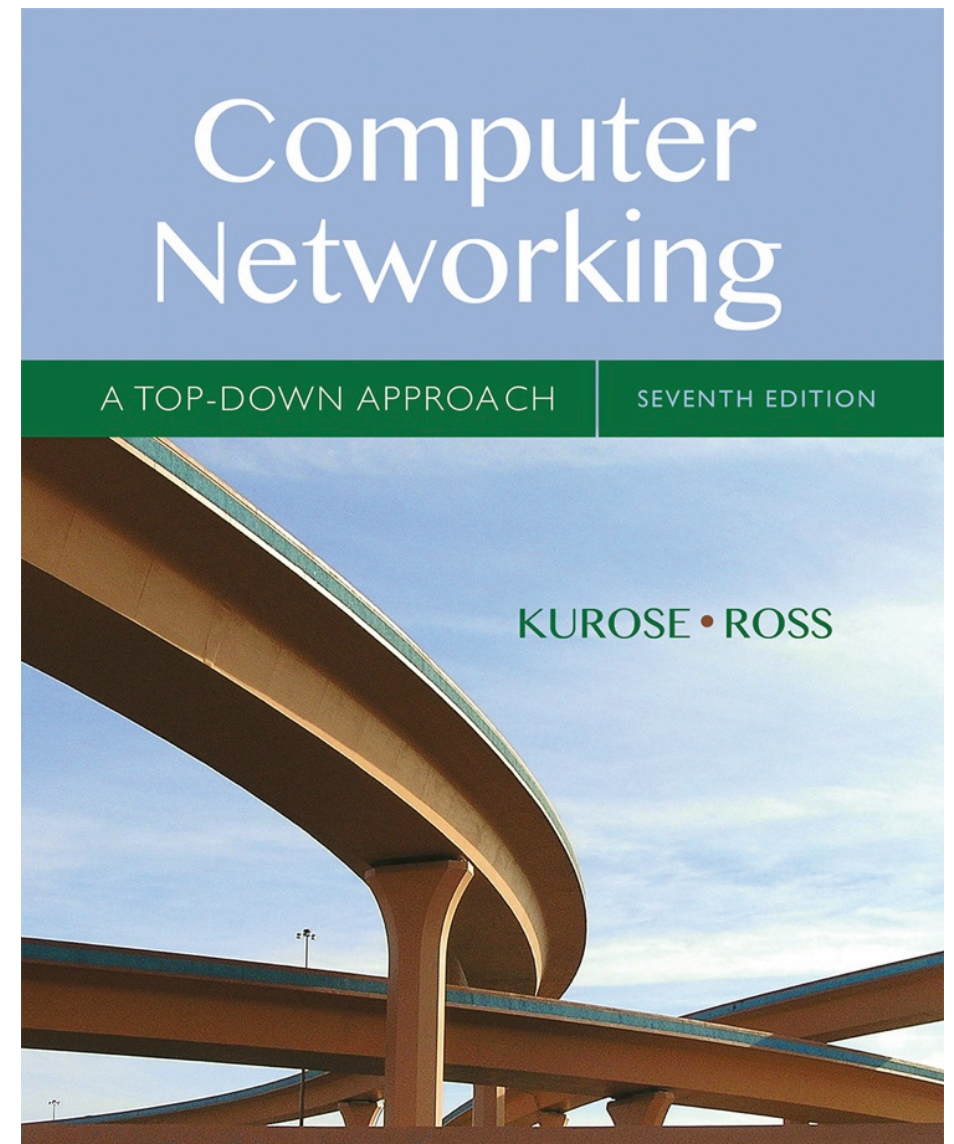
### A note on the use of these Powerpoint slides:

We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

Thanks and enjoy! JFK/KWR

All material copyright 1996-2016  
J.F Kurose and K.W. Ross, All Rights Reserved



## *Computer Networking: A Top Down Approach*

7<sup>th</sup> edition

Jim Kurose, Keith Ross

Pearson/Addison Wesley

April 2016

# P2P Networking

- Need a database of information which peer contains which information.
  - e.g. list of peers having movie-X
    - List could be IP address of peers
  - Information is generally <key, value> pair
- Should it be centralized or distributed
  - Centralized: locating info is easier
  - Distributed: locating info is complex
  - Each have its pros and cons

# Distributed Hash Table (DHT)

- Centralized database
  - <key, value> pair in one central DB, e.g.
    - USN, Name
    - Name, phone number(s)
  - Query the DB with key, get the value(s)
- Challenges with Centralized DB
  - Scaling of DB, Performance,
  - Network congestion
- Solution: Decentralized DB (in P2P version)
  - Each peer holds part of the information.
  - Any peer can query DHT for any info.

# Distributed Hash Table (DHT)

- DHT: a *distributed P2P database*
- Database has **(key, value)** pairs; e.g.:
  - Key: Adhaar number; value: human name
  - Key: movie title; value: IP address
- Distribute the (key, value) pairs
  - over the (millions of peers)
  - A peer may have only key/value info, not actual info
- A peer **queries** DHT with key
  - DHT returns values that match the key
  - Note: It can't query an individual peer
- Peers can also **insert** (key, value) pairs
- Peer can leave too at random

# Use of DHT in P2P Context

- Consider Peers A and B have Linux distro.
- The DHT DB will have: (Linux,  $IP_A$ ), (Linux,  $IP_B$ )
- Consider peer C maintains this key/value pair info.
  - C does not have Linux distro though
  - It only has the info where this distro is available
- Assume that D wants to get this Linux distro.
- D queries the DHT with “Linux” as the key
- DHT decides that C has this key.
- DHT contacts C, and obtains key/value pair.
- This info (Linux,  $IP_A$ ), (Linux,  $IP_B$ ) is given to D
- D contacts either A or B to get the distro.

# Q: how to assign keys to peers?

- Central issue:
  - Can we keep all  $\langle \text{key}, \text{value} \rangle$  pairs in one place
- Distributed implementation
  - How to assign (key, value) pairs to peers?
  - Can these be randomly distributed among peers?
  - Each peer need to know & query every peer
  - How to distribute?
- Basic idea:
  - Convert each key to an integer (Identifier)
  - Assign integer (Identifier) to each peer
  - Put (key,value) pair in the peer **closest** to the key
    - Number of peers doesn't cover full number range

# DHT identifiers


- Assign integer identifier to each peer in range  $[0, 2^n - 1]$  for some  $n$ .
  - Each identifier represented by  $n$  bits.
- Require each key to be an integer in same range
  - Convert given key value (e.g. 17CS52) into integer
- To get integer key, hash original key
  - e.g. key = **hash**("Computer Networks")
  - The reason behind referring it as a *distributed "hash" table*
- *Should each peer know about existence of all other peers?*
  - *Scalability issues?*



# Assign keys to peers

- Rule: assign key to the peer that has the *closest* ID.
- Convention:
  - Closest is the *immediate successor* of the key.
- e.g.,  $n=4$ ; peers: 1, 3, 4, 5, 8, 10, 12, 14;
  - key = 13, then successor peer = 14
  - key = 15, then successor peer = 1
- How does one identify the peer having a key?
  - Each maintains info about assignment
    - Benefit: can locally determine
- How does one know about all existing peers?
  - Issues: not scalable, any changes will cause issues

# Circular DHT (I)



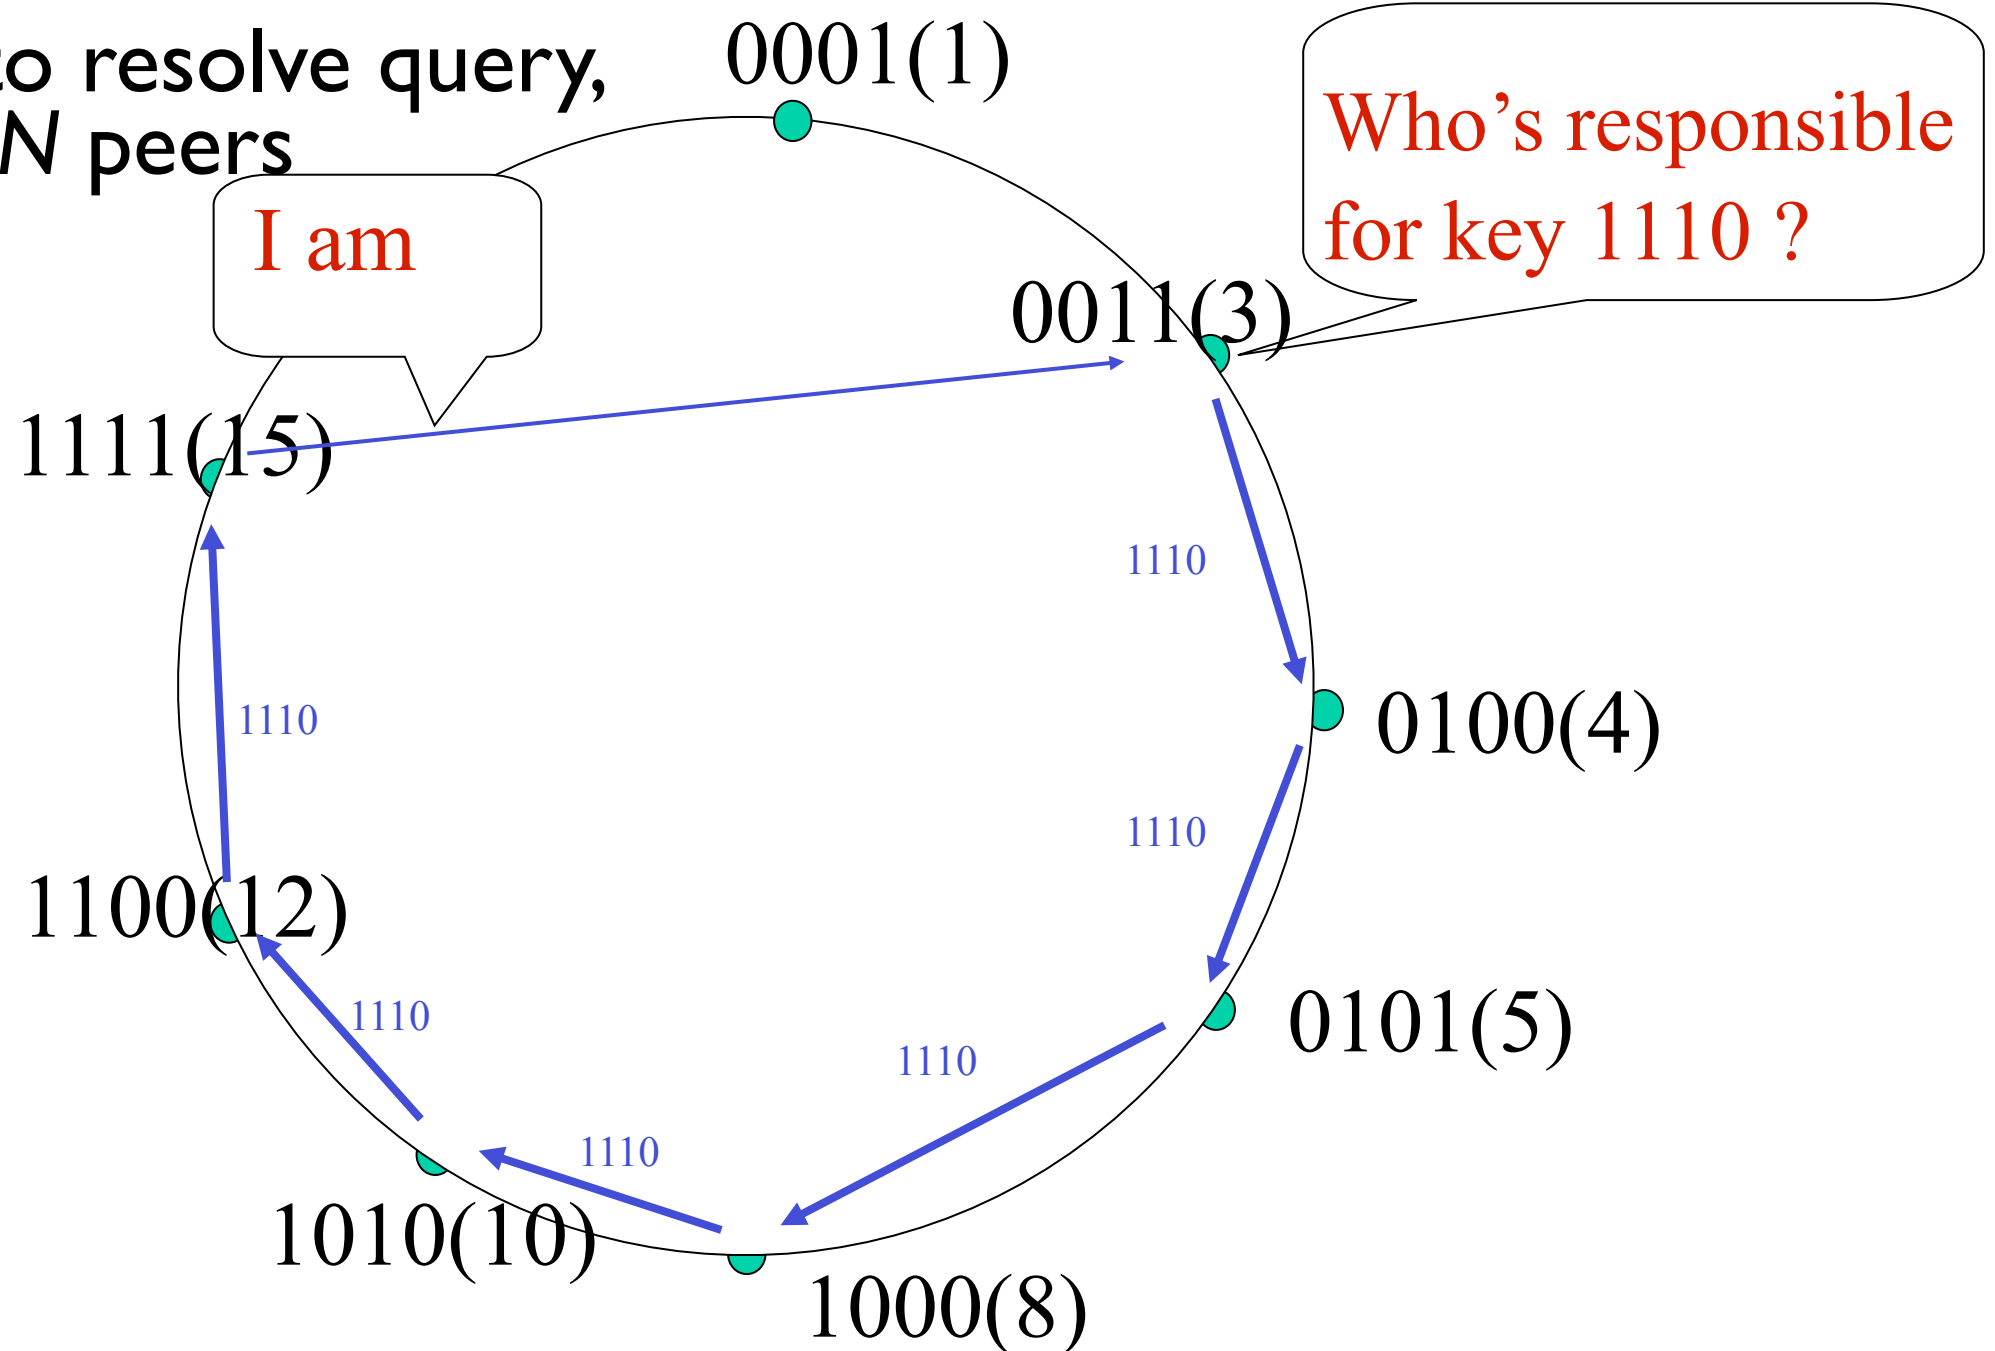
- Each peer *only* aware of immediate successor and predecessor.
- “Overlay network”

# Circular DHT (I)


**Benefit:** Each peer need to know about only 1 peer.

## Challenges:

$O(N)$  messages to resolve query, when there are  $N$  peers



# Circular DHT with shortcuts




- Each peer keeps track of IP addresses of predecessor, successor, and few other short cuts.
- Reduced from 6 to 2 messages.

# Circular DHT with shortcuts

- Trade off between
  - Keeping info about number of peers
  - Number of msgs to communicate for each query
- Two extreme cases
  - Info maintained about all peers
    - communication: 1 msg
  - Info maintained about neighbours
    - communication:  $N/2$  msgs
- Possible to design shortcuts
  - so  $O(\log N)$  neighbours,
  - $O(\log N)$  messages in query

# Peer churn

- *example: peer 5 abruptly leaves*
- Peer 4 detects peer 5 departure;
- makes 8 its immediate successor;
- asks 8 who its immediate successor is;
- makes 8's immediate successor its second successor.
- What if peer 13 wants to join?
  - it only knows peer 1
  - It sends join req to 1
  - This msg keep getting forwarded till it reaches 12
  - 12 knows that it is going to be predecessor of 13
  - 12 informs 13 and accordingly 13 joins the DHT



# Summary

- P2P Distribution
- BitTorrent
- DHT